

Second Mind *by Cognisyn*

Context before composition. Think clearly. Stay yours.

Second Mind turns naturally captured **messages, notes, events, and screenshots** into evidence-linked, revisable context. It then helps people reason and communicate without turning an AI inference into fact or surrendering human agency.

VERIFY IN 60 SECONDS

Demo video (2:28) youtu.be/n53Rs-6uDO4

Browser-ready app codespaces.new/AbbyFitz2-droid/second-mind?quickstart=1 (then: npm start)

Repository github.com/AbbyFitz2-droid/second-mind (public, MIT)

Fastest path Capture a screenshot -> Try a fictional screenshot -> Situation -> People -> Interpretation -> Action.

THE ADVANTAGE: A REASONING ARCHITECTURE, NOT ANOTHER CHAT UI

Invariant evidence. Five lenses keep reported facts fixed while changing the cognitive operation, alternatives, unknowns, and next question. A regression test enforces the contract.

Context before language. Communication Studio uses selected relationship history, goals, preferences, and boundaries across Draft, Reply, Review, Rewrite, Predict, and Compare - without overwriting the user's original.

Inspectable memory. People, situations, commitments, evidence, and interpretations are separate editable objects. AI interpretations never become trusted memory automatically.

Perspective without mind-reading. Perspective Simulation rehearses plausible readings and misunderstanding risk, with evidence and confidence visible.

VERIFIED, REPRODUCIBLE, AND ZERO-COST BY DEFAULT

68 / 68	automated tests pass (npm test)
28 / 28	full-pipeline evaluations pass: intent, schema, and relationship isolation
4,188	ConvoKit conversations in a held-out zero-cost research baseline
0	cloud databases; local judge path needs no private API key

BUILT WITH CODEX + GPT-5.6

Codex /feedback session: 019f5bd7-32bc-77f0-8836-651fd258b040 - one continuous build record covering product decisions, code patches, browser testing, and repairs.

The optional live path uses GPT-5.6 through the OpenAI Responses API with a strict JSON schema. The reproducible judge path uses the same interface and epistemic contract with deterministic local reasoning and paid calls disabled.

Local run: **npm install** | **cp .env.example .env** | **npm start** - no API key needed.

AI should improve the path to an answer - not only the answer.

Most AI products optimise for generating language. Second Mind optimises for constructing a clearer, more inspectable model of reality before language is generated.

"The goal is not to create a conversational AI. The goal is to augment human cognition while preserving human agency."

THE PRODUCT THESIS

People already ask AI to interpret messages, understand difficult relationships, and help with consequential communication. The danger is not only factual hallucination. It is **persuasive certainty built on incomplete context** - a fluent answer that hides the difference between what happened, what was inferred, and what remains unknown.

Second Mind turns that distinction into product architecture. Natural inputs become evidence-linked records. Interpretations remain labelled and revisable. Confidence and alternative explanations remain visible. Changing the cognitive lens cannot silently change the reported facts. Suggestions remain optional, and the user remains the author.

THE DIFFERENCE IS ARCHITECTURAL

GENERAL CONVERSATIONAL AI	SECOND MIND
Starts from the current prompt	Starts from explicit, user-selected context
Memory may remain implicit	People, events, evidence, and commitments are inspectable objects
A fluent interpretation can dominate	Observation, inference, confidence, alternatives, and unknowns remain separate
A new prompt can silently change the frame	Reported evidence stays fixed when the reasoning lens changes
The assistant produces the answer	The person reviews, edits, ignores, and remains the author

THE EDUCATIONAL OUTCOME

Second Mind is designed to reinforce five habits: **What do I know? What am I inferring? How confident am I? What else could explain this? What evidence would change my mind?** The product succeeds when the user becomes more capable of asking those questions without the product.

BUILT WITH CODEX + GPT-5.6

Codex with GPT-5.6 helped turn the thesis into a working system: a schema-versioned context model, local capture pipeline, exact-ID retrieval, five-lens reasoning contract, Communication Studio, browser-tested interactions, and an evaluation suite. The optional live path uses GPT-5.6 through the Responses API with structured output; the zero-cost judge path preserves the same epistemic contract without requiring a private key.

A PERSONAL NOTE

I began with an idea for help during real-world conversations. Building it clarified something deeper: the most valuable AI may not be the system that answers fastest, but the one that helps a person notice how they reached a conclusion.

Thank you to OpenAI for creating the tools, deadline, and space that allowed me to turn that idea into a working, testable product. The process taught me that preserving agency is not a limitation on intelligence. It is a requirement for intelligence people can trust.

Second Mind should leave the user with a clearer model of reality - and still fully the author of the decision.